

```

import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report
from sklearn import tree
from sklearn.tree import DecisionTreeClassifier
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import metrics
%matplotlib inline

file_name = "Dataset.xlsx"
sheet = 'Heart3'

```

```

df = pd.read_excel(io=file_name, sheet_name=sheet)
df.head()

```

	HeartDisease	BMI	Smoking	AlcoholDrinking	Stroke	
0	No	16.60	Yes	No	No	3
1	No	20.34	No	No	Yes	0
2	No	26.58	Yes	No	No	20
3	No	24.21	No	No	No	0
4	No	23.71	No	No	No	28

	MentalHealth	DiffWalking	Sex	AgeCategory	Race	Diabetic	\
0	30	No	Female	55-59	White	Yes	
1	0	No	Female	80 or older	White	No	
2	30	No	Male	65-69	White	Yes	
3	0	No	Female	75-79	White	No	
4	0	Yes	Female	40-44	White	No	

	PhysicalActivity	GenHealth	SleepTime	Asthma	KidneyDisease
0	Yes	Very good	5	Yes	No
1	Yes	Very good	7	No	No
2	Yes	Fair	8	Yes	No
3	No	Good	6	No	No
4	Yes	Very good	8	No	No

```
df=df[['BMI', 'Smoking', 'Sex', 'GenHealth',
'SleepTime', 'SkinCancer']]
df
```

	BMI	Smoking	Sex	GenHealth	SleepTime	SkinCancer
0	16.60	Yes	Female	Very good	5	Yes
1	20.34	No	Female	Very good	7	No
2	26.58	Yes	Male	Fair	8	No
3	24.21	No	Female	Good	6	Yes
4	23.71	No	Female	Very good	8	No
...
319790	27.41	Yes	Male	Fair	6	No
319791	29.84	Yes	Male	Very good	5	No
319792	24.24	No	Female	Good	6	No
319793	32.81	No	Female	Good	12	No
319794	46.56	No	Female	Good	8	No

```
[319795 rows x 6 columns]
```

```
df.isnull().any()
```

```
BMI          False
Smoking      False
Sex          False
GenHealth    False
SleepTime    False
SkinCancer   False
dtype: bool
```

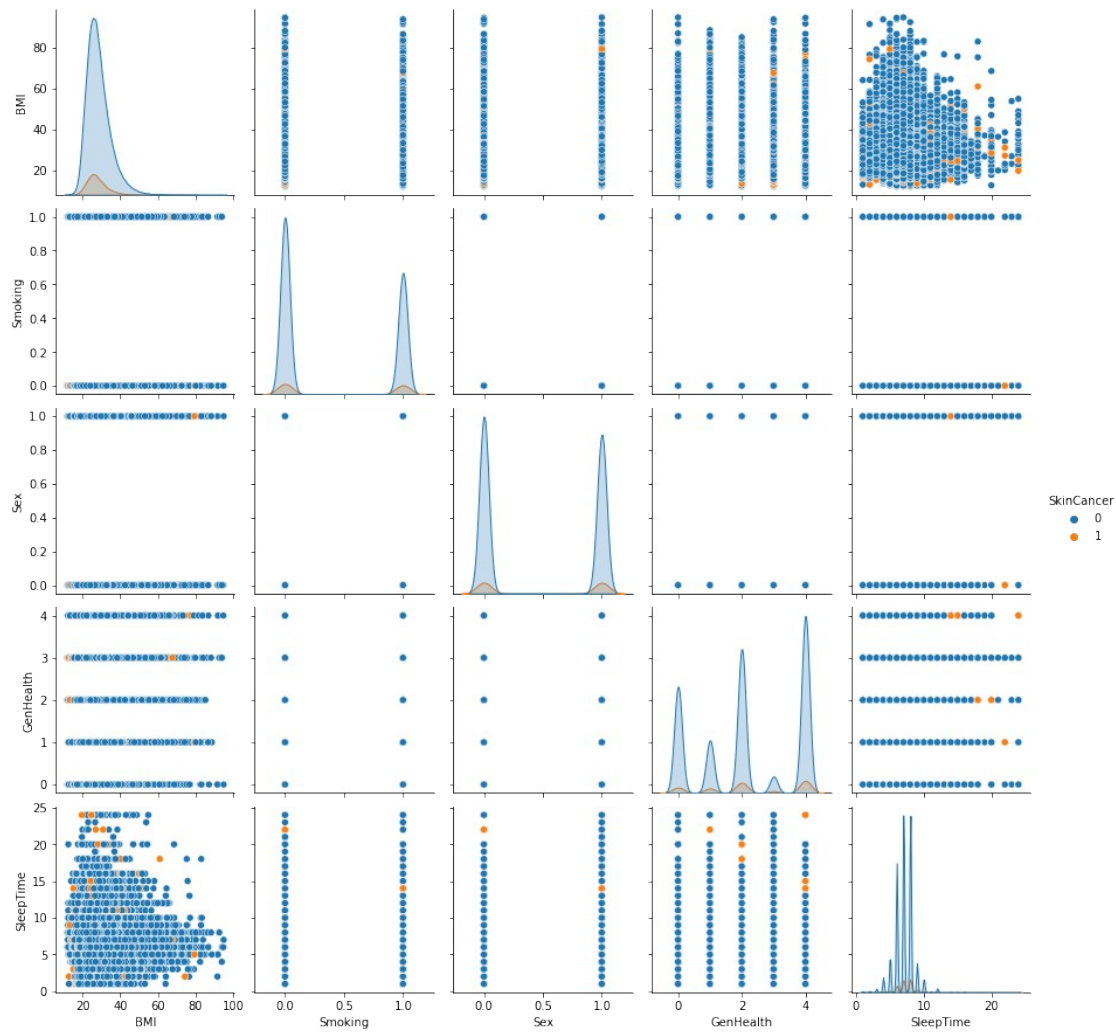
```
from sklearn.preprocessing import LabelEncoder
```

```
enc = LabelEncoder()
```


```
df['Sex'] = enc.fit_transform(df['Sex'])
df['GenHealth'] = enc.fit_transform(df['GenHealth'])
df['Smoking'] = enc.fit_transform(df['Smoking'])
df['SkinCancer'] = enc.fit_transform(df['SkinCancer'])
```

```
sns.pairplot(data=df, hue = 'SkinCancer')
```

```
<seaborn.axisgrid.PairGrid at 0x109e09310>
```



```
corr = df.corr()
plt.figure(figsize=(11,8))
sns.heatmap(corr, cmap="Blues",annot=True)
plt.show()
```



```
target = df['SkinCancer']
df = df.drop('SkinCancer', axis =1)
```

```
X = df
Y = target
```

```
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size =
0.2, random_state = 42)
```

Logistic Regression

```
#logistic regression
```

```
from sklearn import linear_model
```

```
LRG = linear_model.LogisticRegression(random_state = 42,solver =
'liblinear').fit(X_train, y_train)
LRG.score(X_train, y_train)
```

```
0.9068153035538392
```


```
#AUC
```

```
y_pred_proba = LRG.predict_proba(X_test)[::,1]
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred_proba)
auc = metrics.roc_auc_score(y_test, y_pred_proba)
```

```

#create ROC curve
plt.plot(fpr, tpr, label="AUC="+str(auc))
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.legend(loc=4)
plt.show()

```



```

# https://stackoverflow.com/a/61517038/9848043
y_test.reset_index(drop=True)

```

```

0      0
1      1
2      0
3      0
4      1
..
63954   1
63955   0
63956   1
63957   0
63958   0
Name: SkinCancer, Length: 63959, dtype: int64

```

```

# confusion matrix in sklearn
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
# actual values
actual = y_test
# predicted values

```

```

y_preds=LRG.predict(X_test)
predicted = y_preds
# confusion matrix
matrix = confusion_matrix(actual,predicted, labels=[1,0])
print('Confusion matrix : \n',matrix)
# outcome values order in sklearn
tp, fn, fp, tn =
confusion_matrix(actual,predicted,labels=[1,0]).reshape(-1)
print('Outcome values : \n', tp, fn, fp, tn)
# classification report for precision, recall f1-score and accuracy
matrix =
classification_report(actual,predicted,labels=[1,0],zero_division=1)
print('Classification report : \n',matrix)

```

```

Confusion matrix :
[[ 0 5979]
 [ 0 57980]]
Outcome values :
0 5979 0 57980
Classification report :

```


	precision	recall	f1-score	support
1	1.00	0.00	0.00	5979
0	0.91	1.00	0.95	57980
accuracy			0.91	63959
macro avg	0.95	0.50	0.48	63959
weighted avg	0.92	0.91	0.86	63959

```

cm = confusion_matrix(actual, predicted)
plt.figure(figsize=(5,5))
sns.heatmap(data=cm,linewidths=.5, annot=True,square = True, cmap =
'Greens')
plt.ylabel('Actual label')
plt.xlabel('Predicted label')
all_sample_title = 'Accuracy Score:
{:.2f}'.format(accuracy_score(actual,predicted)*100)
plt.title(all_sample_title, size = 15)

Text(0.5, 1.0, 'Accuracy Score: 90.65')

```



Random Forest Classifier

```
from sklearn.ensemble import RandomForestClassifier
```

```
clf = RandomForestClassifier(n_estimators = 100)
```

```
clf.fit(X_train, y_train)
```

```
# performing predictions on the test dataset
```

```
y_pred = clf.predict(X_test)
```

```
# metrics are used to find accuracy or error
```

```
from sklearn import metrics
```

```
# using metrics module for accuracy calculation
```

```
print("ACCURACY OF THE MODEL: ", metrics.accuracy_score(y_test, y_pred))
```

```
ACCURACY OF THE MODEL: 0.8878656639409622
```

```
#AUC
```

```
y_pred_proba = clf.predict_proba(X_test)[::,1]
```


```
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred_proba)
```

```
auc = metrics.roc_auc_score(y_test, y_pred_proba)
```

```
#create ROC curve
```

```
plt.plot(fpr,tpr,label="AUC="+str(auc))
```

```
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.legend(loc=4)
plt.show()
```



```
# confusion matrix in sklearn
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
# actual values
actual = y_test
# predicted values
y_preds=clf.predict(X_test)
predicted = y_preds
# confusion matrix
matrix = confusion_matrix(actual,predicted, labels=[1,0])
print('Confusion matrix : \n',matrix)
# outcome values order in sklearn
tp, fn, fp, tn =
confusion_matrix(actual,predicted,labels=[1,0]).reshape(-1)
print('Outcome values : \n', tp, fn, fp, tn)
# classification report for precision, recall f1-score and accuracy
matrix =
classification_report(actual,predicted,labels=[1,0],zero_division=1)
print('Classification report : \n',matrix)
```

```
Confusion matrix :
[[ 175  5804]
 [ 1368 56612]]
Outcome values :
```



```

175 5804 1368 56612
Classification report :
              precision    recall  f1-score   support

     1         0.11         0.03         0.05         5979
     0         0.91         0.98         0.94        57980


 accuracy          0.89        63959
 macro avg         0.51         0.50         0.49        63959
 weighted avg      0.83         0.89         0.86        63959

```

```

cm = confusion_matrix(actual, predicted)
plt.figure(figsize=(5,5))
sns.heatmap(data=cm,linewidths=.5, annot=True,square = True,  cmap =
'Greens')
plt.ylabel('Actual label')
plt.xlabel('Predicted label')
all_sample_title = 'Accuracy Score:
{:.2f}'.format(accuracy_score(actual,predicted)*100)
plt.title(all_sample_title, size = 15)
Text(0.5, 1.0, 'Accuracy Score: 88.79')

```



K-Nearest Neighbours

```
from sklearn.neighbors import KNeighborsClassifier
```

```
K = []  
training = []  
test = []  
scores = {}
```


```
for k in range(2, 26):  
    clf = KNeighborsClassifier(n_neighbors = k)  
    clf.fit(X_train, y_train)  
  
    training_score = clf.score(X_train, y_train)  
    test_score = clf.score(X_test, y_test)  
    K.append(k)  
  
    training.append(training_score)  
    test.append(test_score)  
    scores[k] = [training_score, test_score]
```

#AUC

```
y_pred_proba = clf.predict_proba(X_test)[::,1]  
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred_proba)  
auc = metrics.roc_auc_score(y_test, y_pred_proba)
```

#create ROC curve

```
plt.plot(fpr,tpr,label="AUC="+str(auc))  
plt.ylabel('True Positive Rate')  
plt.xlabel('False Positive Rate')  
plt.legend(loc=4)  
plt.show()
```



```
for keys, values in scores.items():
    print(keys, ': ', values)

2 : [0.9068231210619303, 0.8979658843946904]
3 : [0.9007020122265826, 0.8838474647821261]
4 : [0.9068231210619303, 0.9032974249128348]
5 : [0.9052947982301162, 0.8989352553979894]
6 : [0.9068113947997937, 0.9057208524210822]
7 : [0.9066003220813333, 0.9042042558514048]
8 : [0.9068426648321581, 0.9060335527447271]
9 : [0.9068270298159758, 0.9055801372754421]
10 : [0.9068622086023859, 0.9063306180521897]
11 : [0.9068348473240669, 0.9060804577932738]
12 : [0.9068270298159758, 0.9064556981816476]
13 : [0.9068543910942948, 0.9063775231007364]
14 : [0.9068113947997937, 0.9064869682140121]
15 : [0.9067918510295658, 0.9064713331978299]
16 : [0.906807486045748, 0.9065182382463766]
17 : [0.9067996685376569, 0.9064713331978299]
18 : [0.9068153035538392, 0.9065182382463766]
19 : [0.9068113947997937, 0.9065026032301944]
20 : [0.9068153035538392, 0.9065182382463766]
21 : [0.906807486045748, 0.9065182382463766]
22 : [0.9068153035538392, 0.9065182382463766]
23 : [0.9068153035538392, 0.9065182382463766]
24 : [0.9068153035538392, 0.9065182382463766]
25 : [0.9068153035538392, 0.9065182382463766]
```

```

# confusion matrix in sklearn
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
# actual values
actual = y_test
# predicted values
y_preds=clf.predict(X_test)
predicted = y_preds
# confusion matrix
matrix = confusion_matrix(actual,predicted, labels=[1,0])
print('Confusion matrix : \n',matrix)
# outcome values order in sklearn
tp, fn, fp, tn =
confusion_matrix(actual,predicted,labels=[1,0]).reshape(-1)
print('Outcome values : \n', tp, fn, fp, tn)
# classification report for precision, recall f1-score and accuracy
matrix =
classification_report(actual,predicted,labels=[1,0],zero_division=1)
print('Classification report : \n',matrix)

```

```

Confusion matrix :
[[ 0 5979]
 [ 0 57980]]
Outcome values :
0 5979 0 57980
Classification report :

```


	precision	recall	f1-score	support
1	1.00	0.00	0.00	5979
0	0.91	1.00	0.95	57980
accuracy			0.91	63959
macro avg	0.95	0.50	0.48	63959
weighted avg	0.92	0.91	0.86	63959

```

cm = confusion_matrix(actual, predicted)
plt.figure(figsize=(5,5))
sns.heatmap(data=cm,linewidths=.5, annot=True,square = True, cmap =
'Greens')
plt.ylabel('Actual label')
plt.xlabel('Predicted label')
all_sample_title = 'Accuracy Score:
{:.2f}'.format(accuracy_score(actual,predicted)*100)
plt.title(all_sample_title, size = 15)

Text(0.5, 1.0, 'Accuracy Score: 90.65')

```




```
ax = sns.stripplot(K, training);
ax.set(xlabel='values of k', ylabel='Training Score')
```

```
plt.show()
```


/Users/joyanta/miniforge3/envs/tensorflow_3/lib/python3.8/site-packages/seaborn/_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```




```
ax = sns.stripplot(K, test);  
ax.set(xlabel = 'values of k', ylabel = 'Test Score')  
plt.show()
```

/Users/joyanta/miniforge3/envs/tensorflow_3/lib/python3.8/site-packages/seaborn/_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn(



```
plt.scatter(K, training, color='k')
plt.scatter(K, test, color='g')
plt.show()
```



Naive Bayes

```
from sklearn.naive_bayes import GaussianNB
gnb = GaussianNB()
gnb.fit(X_train, y_train)
```


```
# making predictions on the testing set
y_pred = gnb.predict(X_test)

# comparing actual response values (y_test) with predicted response
values (y_pred)
from sklearn import metrics
print("Gaussian Naive Bayes model accuracy(in %):",
metrics.accuracy_score(y_test, y_pred)*100)
```

Gaussian Naive Bayes model accuracy(in %): 90.65182382463766

```
#AUC
y_pred_proba = gnb.predict_proba(X_test)[::,1]
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred_proba)
auc = metrics.roc_auc_score(y_test, y_pred_proba)
```

```
#create ROC curve
plt.plot(fpr,tpr,label="AUC="+str(auc))
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.legend(loc=4)
plt.show()
```



```
# confusion matrix in sklearn
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
# actual values
actual = y_test
# predicted values
y_preds=gnb.predict(X_test)
```



```

predicted = y_preds
# confusion matrix
matrix = confusion_matrix(actual,predicted, labels=[1,0])
print('Confusion matrix : \n',matrix)
# outcome values order in sklearn
tp, fn, fp, tn =
confusion_matrix(actual,predicted,labels=[1,0]).reshape(-1)
print('Outcome values : \n', tp, fn, fp, tn)
# classification report for precision, recall f1-score and accuracy
matrix =
classification_report(actual,predicted,labels=[1,0],zero_division=1)
print('Classification report : \n',matrix)

```

```

Confusion matrix :
[[ 0 5979]
 [ 0 57980]]
Outcome values :
0 5979 0 57980
Classification report :

```


	precision	recall	f1-score	support
1	1.00	0.00	0.00	5979
0	0.91	1.00	0.95	57980
accuracy			0.91	63959
macro avg	0.95	0.50	0.48	63959
weighted avg	0.92	0.91	0.86	63959

```

cm = confusion_matrix(actual, predicted)
plt.figure(figsize=(5,5))
sns.heatmap(data=cm,linewidths=.5, annot=True,square = True, cmap =
'Greens')
plt.ylabel('Actual label')
plt.xlabel('Predicted label')
all_sample_title = 'Accuracy Score:
{:.2f}'.format(accuracy_score(actual,predicted)*100)
plt.title(all_sample_title, size = 15)

Text(0.5, 1.0, 'Accuracy Score: 90.65')

```



Support Vector Machine

```
# import support vector classifier
from sklearn.svm import SVC
from sklearn.svm import LinearSVC
# clf = SVC(kernel='linear') #it takes too much time to train up
clf = LinearSVC(random_state=42, tol=1e-5)
# fitting x samples and y classes
clf.fit(X_train, y_train)
```

```
y_pred = clf.predict(X_test)
```

```
print("SVM model accuracy(in %):", metrics.accuracy_score(y_test,
y_pred)*100)
```

```
SVM model accuracy(in %): 90.65182382463766
```


```
/Users/joyanta/miniforge3/envs/tensorflow_3/lib/python3.8/site-
packages/sklearn/svm/_base.py:1206: ConvergenceWarning: Liblinear
failed to converge, increase the number of iterations.
  warnings.warn(
```

#AUC

```
y_pred_proba = clf.predict(X_test)
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred_proba)
auc = metrics.roc_auc_score(y_test, y_pred_proba)
```

#create ROC curve

```
plt.plot(fpr, tpr, label="AUC="+str(auc))
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.legend(loc=4)
plt.show()
```



```
# confusion matrix in sklearn
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
# actual values
actual = y_test
# predicted values
y_preds=clf.predict(X_test)
predicted = y_preds
# confusion matrix
matrix = confusion_matrix(actual,predicted, labels=[1,0])
print('Confusion matrix : \n',matrix)
# outcome values order in sklearn
tp, fn, fp, tn =
confusion_matrix(actual,predicted,labels=[1,0]).reshape(-1)
print('Outcome values : \n', tp, fn, fp, tn)
# classification report for precision, recall f1-score and accuracy
matrix =
classification_report(actual,predicted,labels=[1,0],zero_division=1)
print('Classification report : \n',matrix)

Confusion matrix :
[[ 0 5979]
 [ 0 57980]]
```

Outcome values :

0 5979 0 57980

Classification report :

	precision	recall	f1-score	support
1	1.00	0.00	0.00	5979
0	0.91	1.00	0.95	57980
accuracy			0.91	63959
macro avg	0.95	0.50	0.48	63959
weighted avg	0.92	0.91	0.86	63959

```
cm = confusion_matrix(actual, predicted)
plt.figure(figsize=(5,5))
sns.heatmap(data=cm,linewidths=.5, annot=True,square = True, cmap =
'Greens')
plt.ylabel('Actual label')
plt.xlabel('Predicted label')
all_sample_title = 'Accuracy Score:
{:.2f}'.format(accuracy_score(actual,predicted)*100)
plt.title(all_sample_title, size = 15)
Text(0.5, 1.0, 'Accuracy Score: 90.65')
```

